

Marking Scheme - PA3

110 MARKS Undergraduate

130 MARKS Graduate

Q1 - MNIST Download (2 Marks)

- 1 Mark: Dataset downloaded with torchvision
- 1 Mark: Transforms applied & dataset size printed

Q2 - Splits (2 Marks)

- 1 Mark: Train/Val/Test splits are correct (50k/10k/10k)
- 1 Mark: DataLoader sizes printed

Q3 - Visualisation (2 Marks)

- 1 Mark: visualise_samples implemented with correct arguments
- 1 Mark: At least 2 samples displayed

Q4 - Forward Pass in NumPy (17 Marks)

1. **He Initialization** - 5 Marks
 - 3 Marks: Correct formula
 - 2 Marks: Bias set to zeros
2. **Weight Initialization** - 2 Marks (Correct dimensions per layer)
3. **Activations** - 5 Marks
 - 2 Marks: ReLU correct
 - 3 Marks: Softmax correct (with stability)
4. **Forward Pass** - 5 Marks
 - 3 Marks: Correct probability outputs
 - 2 Marks: Predicted class printed

Q5 - PyTorch Neural Networks (62 Marks)

1. **Count Parameters** - 2 Marks (Binary)
2. **FCNN (>200k params)** - 10 Marks
 - 4 Marks: Correct dimensions / architecture
 - 2 Marks: ReLU activation used
 - 2 Marks: Param count >200k shown
 - 2 Marks: No convolutions used

3. Hyperparameters - 5 Marks

- 2 Marks: Correct loss function
- 2 Marks: Optimizer + LR set
- 1 Mark: Explanation in markdown

4. Training s Evaluation Functions - 20 Marks

- 7.5 Marks: train function implemented
- 7.5 Marks: evaluate function implemented
- 5 Marks: Average loss accuracy correctly computed
- -2 Marks: Per wrong explanation of `zero_grad()`, `backward()`, `step()`

5. Training Loop - 15 Marks

- 5 Marks: Training/Validation loss C accuracy printed for each epoch
- 5 Marks: Best model saved properly
- 5 Marks: Curves plotted and labelled

6. Evaluation on Test Set - 10 Marks

- 3 Marks: Accuracy $\geq 95\%$
- 3 Marks: F1 C Recall $\geq 95\%$
- 2 Marks: Classification report + confusion matrix
- 2 Marks: ROC curve + incorrect samples visualised

7. OOD Evaluation - 15 Marks

- 13 Marks: Image preprocessing + visualization + Forward Pass
- 1 Marks: Predictions shown
- 1 Mark: Comment on performance

Q6 - Data Pruning s Smaller Model (25 Marks)

1. Smaller FCNN (<110k params) - 5 Marks

- 3 Marks: Correct architecture
- 2 Marks: Param count printed

2. Skip

3. Hyperparameters - 2 Marks (Binary)

4. Bottom-k Dataloader - 5 Marks (Subset used correctly)

5. **Training Loop** - 5 Marks (Metrics printed, works end-to-end)

6. **Evaluation** - 8 Marks

- 3.5 Marks: Accuracy $\geq 95\%$
- 3.5 Marks: F1 C Recall $\geq 95\%$
- 1 Marks: Confusion matrix + performance comparison

- -5 if average loss not computed correctly

- -5 if accuracy computed incorrectly

Graduate Question MS:

Per Activation Function 1 Mark -

3 Marks (Binary)

Backward Pass - 17 Marks:

Correct Derivative of Loss Function - BCE -

2 Marks

Correct Shapes of dW , dB , L , m -

4 Marks (1 each)

Correct For Loop for Backpropogation:

Goes from L to 1 -

-5 if incorrect

Correctness of dW

3 Marks

Correctness of dB

3 Marks

Correctness of dZ (for next iteration) - 5 Marks